

# **CHUMP**

## **A Mechanized Verification of Curricule**

**Justin Frank and Pierce Darragh**

**CMSC 838L**

**May 1, 2024**

# Curtailed Curricule Curriculum

- A type system for safe intermittent computing
- Generalization over various implementations of checkpointing
- Support for both just-in-time and atomic intermittent execution
- Inference of checkpointing from (non-)idempotence annotations
- Built on top of Rust

# Our Motivation

- Curricle is cool
- Mechanization is *also* cool
- Want a mechanized semantics of intermittent systems
  - Proving correctness on intermittent semantics directly is *hard*
  - Want to write proofs about continuous execution that translate to intermittent behavior

# Mechanizing Curricula

- Extended paper has...
  - ~10 pages of typing rules, inference rules, and related material
  - ~10 pages of formal semantics
  - ~30 pages of proofs
- It would be a shame if something were to... *happen* while implementing it all
- Solution: mechanization

# Project Outline

- Simplify language and feature set
  - But add a little something new, as a treat
- Decide on a meaning of "correctness"
- Prove it, using adjusted versions of theorems and relations from the paper

# Language and Features

# Some Simplifications

(Because mechanization is hard)

- Simplified language
  - No structs
  - No functions
- Simplified semantics
  - Only undo checkpointing
  - DINO-style checkpoints
- Simplified type system
  - Separate Curricule from base types
- Simplified machine state
  - No volatile memory


# Additional Features

- Support for output in non-JIT intermittent execution
- Unbounded loops with arbitrarily nested breaks

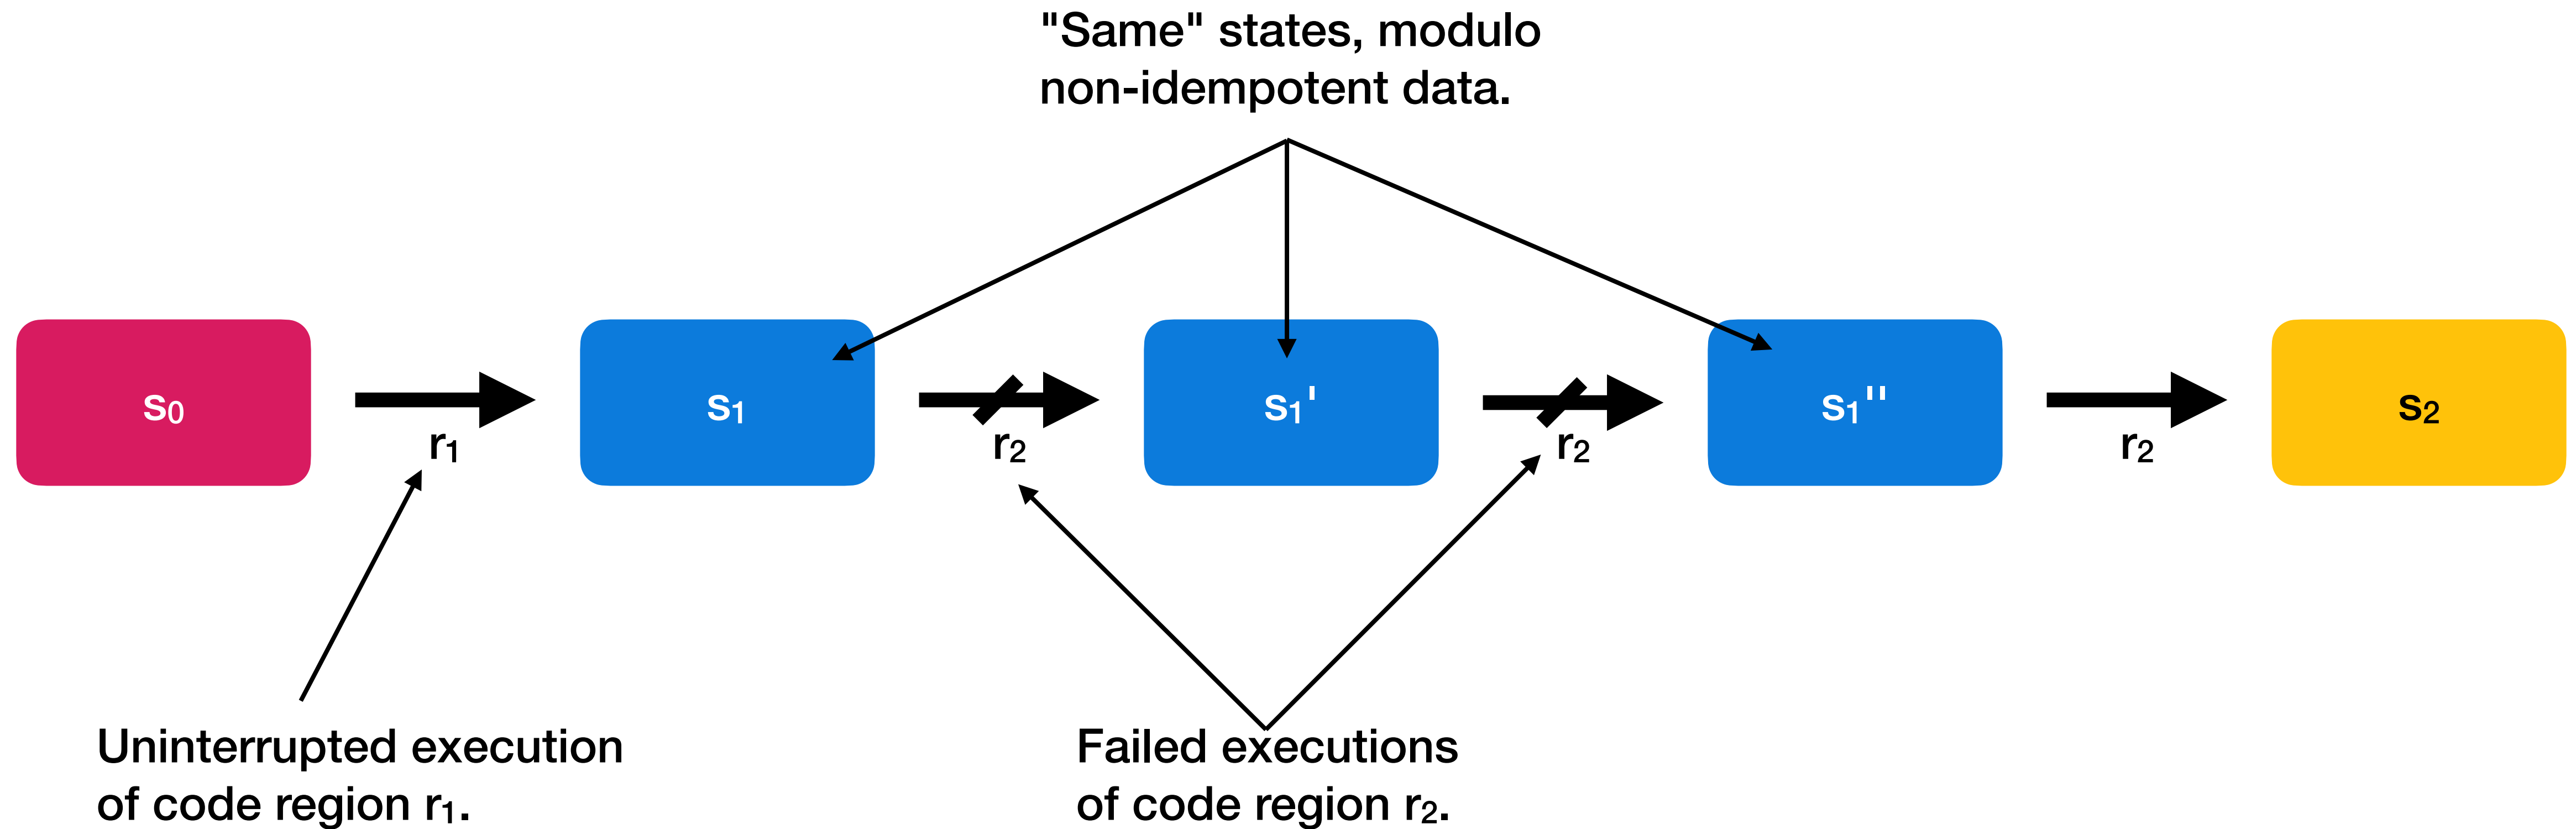


# Correctness

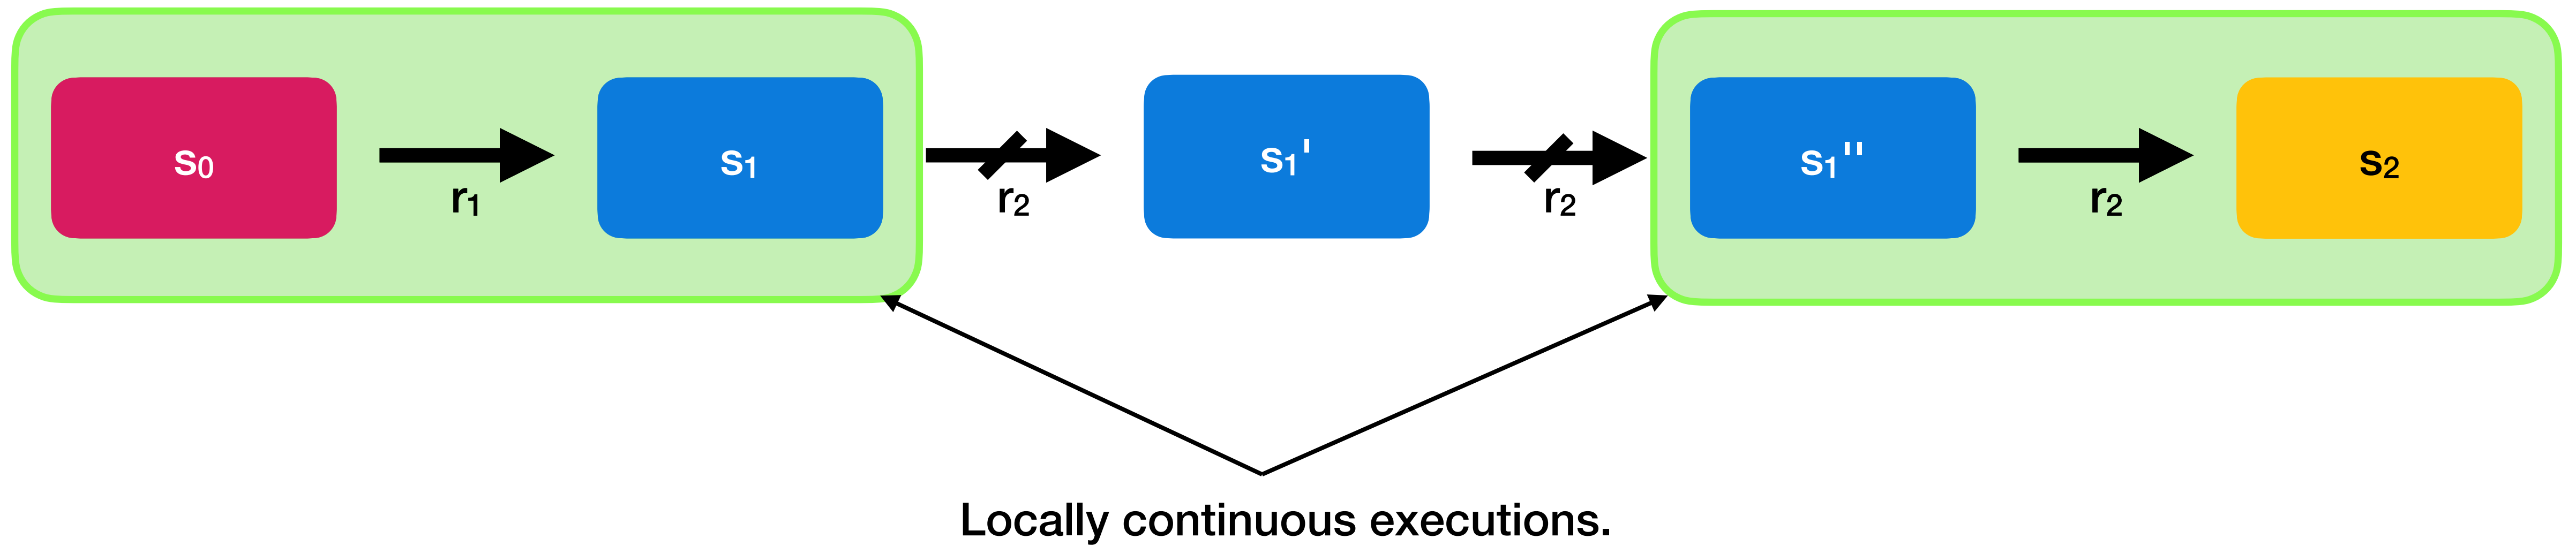
# Correctness



# Correctness




# Correctness



# Correctness



Intermittent execution.

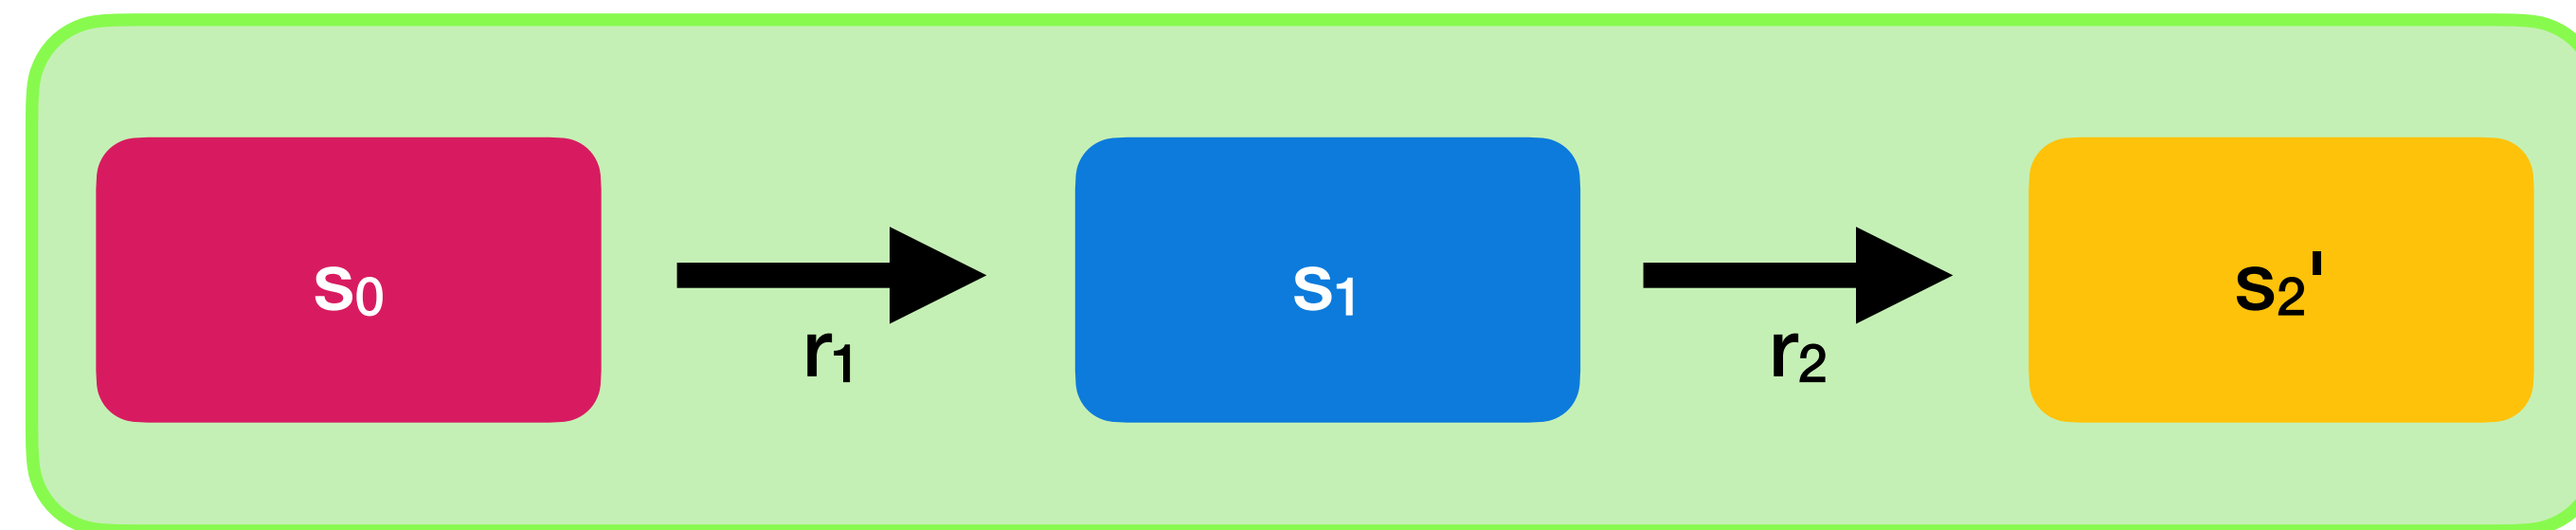


Simulation of continuous execution.

# Correctness



Intermittent execution.

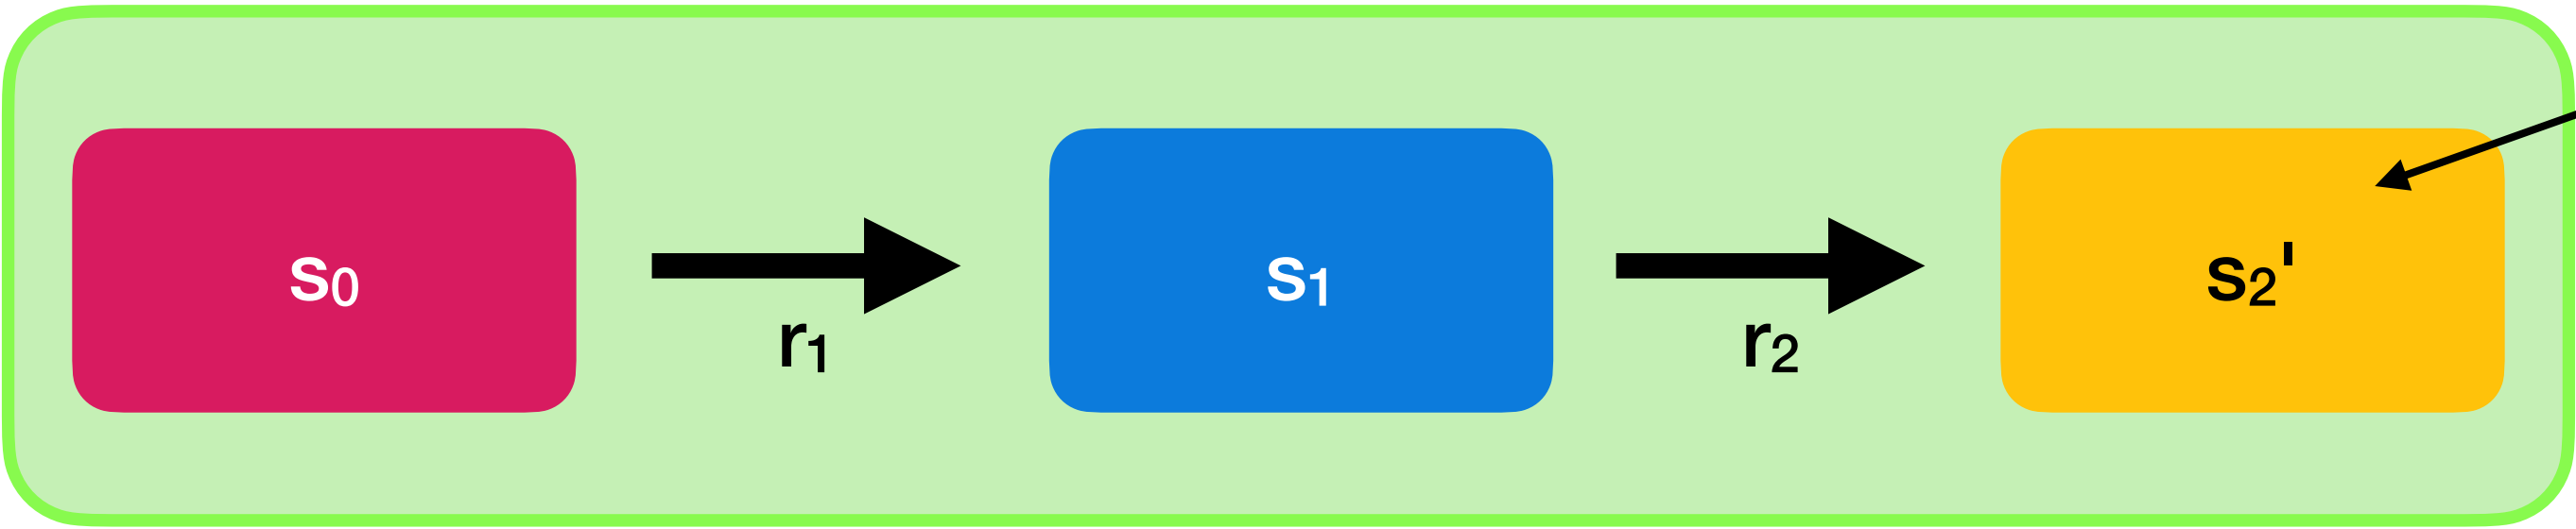


Simulation of continuous execution.

# Correctness



Intermittent execution.



Simulation of continuous execution.

"Same" states, modulo  
non-idempotent data.

# Correctness

- Curricule has a two-part notion of correctness:
  1. Non-interference: power failures do not adversely affect end result
  2. Intermittent exec. w/ no power failures is equivalent to continuous exec.
- We changed it a little bit:
  - $\forall$  exec.,  $\exists$  reset-free exec. equivalent to the original exec.



# Proof Approach

# Machine Model

- Extended CESK with checkpoints, making CESKP
  - Control, Environment, Store, Kontinuation, checkPoint
- CESK is a semantic model both team members are comfortable with
- Regularly used in mechanized formalizations (so should be "easy" to adapt)

```
Inductive kont          := { ... }.
Definition addr         := positive.
Inductive val          := { ... }.
Definition env          := list.
Inductive c            := { ... }.
Definition checkpoint : Type := kont * list (addr * val).
Definition CESKP      : Type := c * env addr * store val * kont * checkpoint.
```

# The Outside World

- Input and output are modeled as external constructs
  - i.e., separate from the CESKP machine state
- Resets are considered IO events

```
Variant io_event :=  
| io_in  (i : Z)  
| io_out (i : Z)  
| io_check  
| io_reset.  
Definition io_log      := list io_event.  
Definition state : Type := CESKP * io_log.
```

# Separating the Type Systems

- Curricule's type system is all-in-one
- We separate base types from Curricule's type qualifiers

$$\frac{
 \begin{array}{l}
 \Pi(p) = l_p \quad \Gamma, \Pi, \text{Wts} \vdash_e^{\text{atom}} e : t @ q_e \Rightarrow \Gamma' \quad \Gamma', \Pi, \text{Wts} \vdash_{\text{elft}}^{\text{atom}} p : @q_l \Rightarrow \Gamma'_l \\
 \langle qId_e, qIO_e \rangle = q_e \sqcup \langle qId_{pc}, qIO_{pc} \rangle \sqcup q_l \quad \Gamma'_l(l_p) = t @ \langle qId_p, qIO_p \rangle \\
 qId_e \sqsubseteq qId_p \downarrow \quad qId_p \uparrow^{qAcc} \leq_a \text{Wt} \quad \text{UpdatePoint}(\Pi, p, e : t) = \Pi'
 \end{array}
 }{
 \Gamma, \Pi, \text{Wts}, \langle qId_{pc}, qId_{pc} \rangle \vdash_c^{\text{atom}} p := e \Rightarrow \Gamma'_l[\Pi'(p) : \langle qId_p, qIO_e \rangle], \Pi', (\text{Wts}, \Pi'(p))
 } \text{ASSGNSINGLE}$$

# Energy Model

- Curricule uses logical energy
  - Energy is encoded into the proofs as a decreasing natural number
  - $\text{Energy} = 0 \implies \text{power failure}$
- We use a nondeterministic semantic model
  - Any state can transition to a reset instead of its usual next step

# Side Quests

- Proofs of *progress* and *preservation* in the base type system
  - ~800 lines of Rocq
- (Most of) a formalization of lattices for the Curricule types
  - ~100 lines of Rocq

# Side Quests

- Proofs of *progress* and *preservation* in the base type system
  - ~500 lines of Rocq
  - Not necessary
- (Most of) a formalization of lattices for the Curricle types
  - ~100 lines of Rocq
  - Not necessary

# Side Quests

## 15.2 Proving in Coq

While Coq proofs can be checked a posteriori in batch mode, they are developed interactively using a number of tactics as elementary proof steps. The sequence of tactics used constitutes the proof script. Building such scripts is surprisingly addictive, in a videogame kind of way, but reading them is notoriously difficult. We were initially concerned that adapting and reusing proof scripts when specifications change could be difficult, forcing many proofs to be rewritten from scratch. In practice, careful decomposition of the proofs in separate lemmas enabled us to reuse large parts of the development even in the face of major changes in the semantics, such as switching from the “mixed-step” semantics described in [57] to the small-step transition semantics described in this paper.

— "A formally verified compiler back-end", Xavier Leroy, 2009.



# What Remains

# Current Status

- Still implementing Curricule-based type system
- Getting to this point has taken *way* more work than we'd thought
- Plan to finish basic proofs before the end of the semester

# Future Work

- Add missing language features (structs, functions, ...?)
- Implement a DSL (maybe in Racket, just for fun)

**End.**



# Why "CHUMP"?

- IMP is a basic imperative language used for modeling semantics
- CHeckpoints + Undo + iMP = CHUMP